

Real-Time Memory Management and Compaction with Bounded Response Times

Hannes Payer

University of Salzburg

July 11, 2007

Introduction

Project Goals

Memory Model

Design

Implementations

General Page Concept

Virtual Memory

Physical Memory


Motivation

- traditional memory management is typically non-deterministic:
 - unpredictable response times
 - worst-case response times are usually unbounded
 - memory fragmentation
- for WCET analysis it is important that all the running code (application, libraries, and operating system) is predictable

Motivation

- response time and memory compaction should be independent of the memory context
- integration in a cooperative system:
 - no preemption
 - incremental behaviour of each operation

Motivation



Source: *M. Masmano, I. Ripoll, and A. Crespo. Dynamic storage allocation for real-time embedded systems, 2003.*

Goals and Requirements

memory operations malloc, free, dereference

- bounded response time (predictability)
- fast response time


fragmentation

- eliminated by compaction
- distribute compaction workload (deterministic timing behaviour)
- predictable timing solution for reference updates required


Compaction & Abstract vs. Concrete Address Space

- reference updates degrades predictability
- compaction with reference updates can not be done fast and bounded?!
- \Rightarrow memory operations operate on abstract addresses to be independent of compaction operations
- abstract address refers to concrete addresses
- abstract address never changes
- concrete addresses may change due to compaction

Proxy - abstract to concrete address mapping



Proxy - abstract to concrete address mapping



Memory Operations API

- `void **malloc(size)`: allocates memory and returns an abstract address that points to the concrete address space
- `void free(abstract address)`: frees the concrete address space that belongs to the abstract address
- `void *dereference(abstract address, offset)`: returns an address of the concrete address space that corresponds to the abstract address + offset

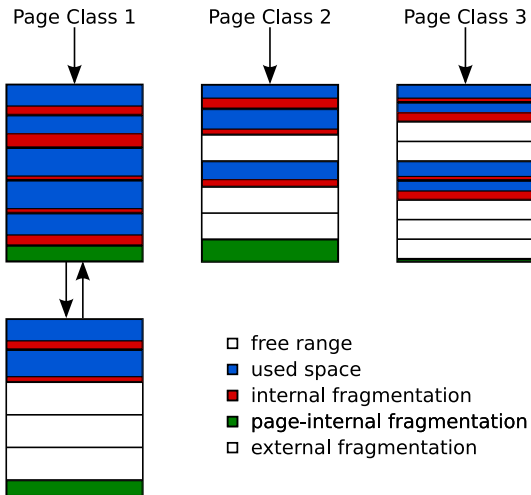
Pages (Metronome)

- concrete address space is divided into pages of equal size
- each page itself is divided into fixed-sized page blocks
- there exist n predefined page block sizes, which results in n different size classes
- allocation operations are handled by the smallest size class that fits the allocation request
- segregated free lists
- redistribute unused pages to other size classes

Fragmentation (Metronome)

- internal fragmentation (unused space at the end of a page block)
- page-internal fragmentation (unused space at the end of a page)
- external fragmentation (unused page blocks)

Page Classes & Fragmentation (Metronome)



Compaction

Invariant 1: There exists at most one page in a page class, which is not full.

Invariant 2: The not-full-page has to be the last element of its page class list.

Rule 1: If an element of a full page gets freed and there exists no not-full-page, then this page is the new not-full-page of the page class and it is moved to the end of the page class list.

Rule 2: If an element of a full page gets freed and there exists a not-full-page, then one element of the not-full-page has to be moved to this page.

Compaction




Figure: fragmented pages

Compaction




Figure: free one memory range

Compaction




Figure: move memory range of the last page to the affected page

Page Management Internals

administration information

- concrete address space
- free entries counter
- next/previous pointer to the next/previous page of the page class (doubly-circularly-linked list)
- free page block list
- used page block list
- reference to the page class

Page Management Internals

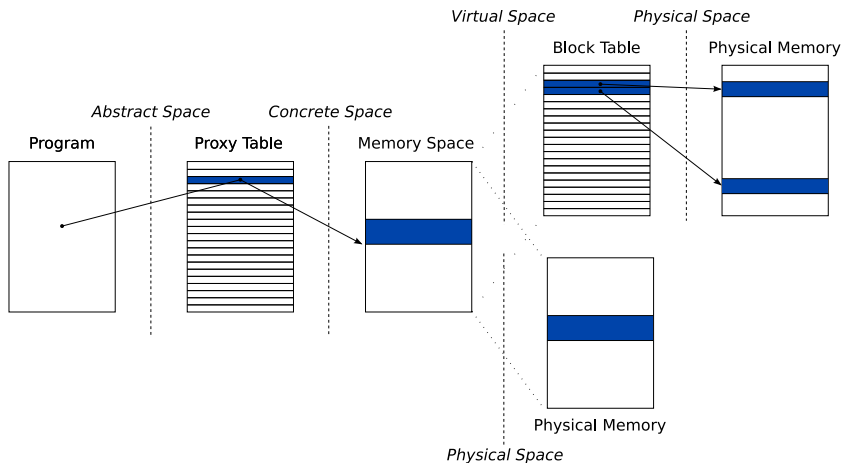
free page block list

- singly-linked list (stack) for free page blocks of a page
- special list construct
- no list initialization (no free element concatenation)
- build up free list incrementally
- 16 bits needed for link pointer

used page block list

- needed for compaction
- doubly-circularly-linked list (32 bits for link pointers)
- 2-dimensional bitmap (1 bit for each page block)

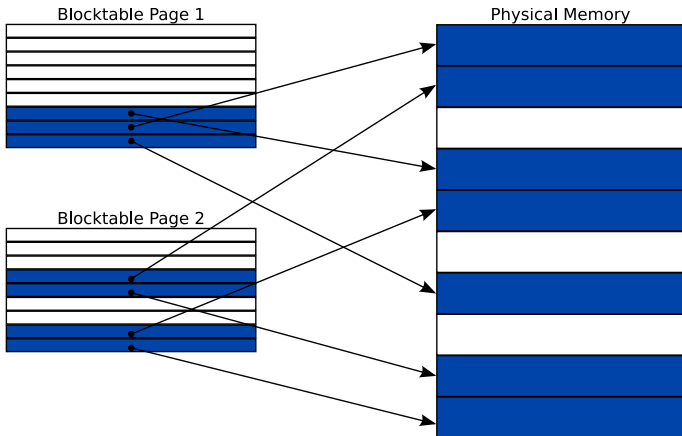
Overview - 2 different implementations



Virtual Memory

- the virtual memory is divided into blocks of equal size
- the physical memory is divided into block frames of equal size
- allocated memory objects are contiguous in virtual space, but arbitrarily distributed in physical space
- compaction happens in virtual space - block frames are never moved in physical space
- constant compaction speedup
- every page contains a block table that implements the virtual-block to physical-block-frame mapping

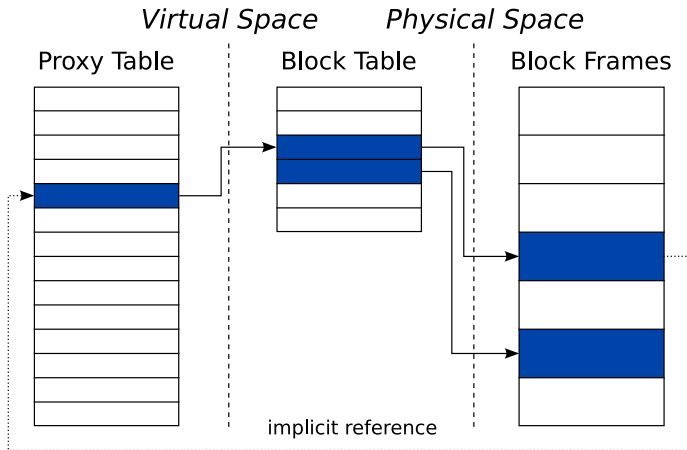
Block Table



Proxy Name-Generation

- allocated memory is never moved physically
- # block frames = # proxy entries
- proxy entry assignment in constant time
- given the concrete address of a concrete memory object, the abstract address can be determined implicitly

Proxy Implicit Reference



void **malloc(size)

```
page = getPageOfPageClass(size);
freeIndexInPage = getFreeIndexOfPage(page);
for(i=0; i<numberOfBlockFrames; i++){
    blockframe = getBlockFrameFromFreeList();
    page->bf[index+i] = blockframe;
}
addUsedIndexToPage(page, index);
proxyIndex = getIndexOfBlockFrame(page->bf[index]);
return getProxyEntry(proxyIndex, &page->bf[index]);
```

void free(abstractAddress)

```
concreteAddress = dereference(abstractAddress);
page = getPageToConcreteAddress(concreteAddress);
index = indexOfRangeInPage(page, concreteAddress);
removeUsedIndexFromPage(page, index);
for(i=0; i<numberOfBlockFrames; i++){
    addBlockFrameToFreeList(page->bf[indexOfBf+i]);
}
addFreeIndexToPage(page, index);

if(page->entries == 0){
    removePageFromPageClass(page);
}
else{
    doCompaction(page);
}
```


`void *dereference(abstractAddress, index)`

```
concreteAddress = dereference(abstractAddress);  
return getWordReference(concreteAddress, index);
```

Conclusion I

- malloc: $\Theta(\frac{n}{Block\ Frame\ Size})$ timing behaviour (constant in page class size)
- free: $\Theta(\frac{n}{Block\ Frame\ Size})$ timing behaviour (constant in page class size)
 - $\frac{n}{Block\ Frame\ Size}$ block table free operations
 - $\frac{n}{Block\ Frame\ Size}$ compaction operations (if compaction occurs)
- dereference: $\Theta(1)$ timing behaviour
- the timing behaviour of all memory operations is independent of the global memory state

Overview - 2 different implementations




Physical Memory

- the concrete address space is the physical space \Rightarrow physical space is divided into pages of equal size
- allocated memory objects are contiguous in physical space
- compaction leads to movements in physical space

Proxy Name-Generation

- physical space locations change due to compaction operations
- free list over the proxy table (constant time)
- beginning address of the memory object is not related to the proxy table entry position
- \Rightarrow memory objects need an explicit reference back to the corresponding proxy entry

Proxy Explicit Reference



void **malloc(size)

```
page = getPageOfPageClass(size);  
freeIndexInPage = getFreeIndexOfPage(page);  
addUsedIndexToPage(page, index);  
proxyIndex = getIndexOfBlockFrame(page->bf[index]);  
return getProxyEntry(proxyIndex, &page->bf[index]);
```

void free(abstractAddress)

```
concreteAddress = dereference(abstractAddress);  
page = getPageToConcreteAddress(concreteAddress);  
index = indexOfRangeInPage(page, concreteAddress);  
removeUsedIndexFromPage(page, index);  
addFreeIndexToPage(page, index);  
  
if(page->entries == 0){  
    removePageFromPageClass(page);  
}  
else{  
    doCompaction(page);  
}
```

```
void *dereference(abstractAddress, index)
```

```
just use the *(*abstractAddress + index)  
explicit dereference method is not necessary
```

Conclusion II

- malloc: $\Theta(1)$ timing behaviour
- free: $\mathcal{O}(n)$ timing behaviour (constant in range size)
 - $\Theta(1)$ free operation
 - n compaction operations $\Rightarrow \Theta(n)$
- dereference: $\Theta(1)$ timing behaviour
- the timing behaviour of all memory operations is independent of the global memory state